Graphic Lasso: Clustering accuracy for precision matrix model

Jiaxin Hu

February 18, 2021

1 With convex penalty L_1 norm

The precision model is stated as

$$\mathbb{E}[S^k] = \Omega^k = \sum_{l=1}^{r} u_{kl}\Theta^l, \quad k \in [K].$$

Consider the following penalized optimization problem

$$\max_{\boldsymbol{U},\Theta^l} \mathcal{L}_S(\boldsymbol{U}, \Theta^l) = -\sum_{k=1}^{K} \operatorname{tr}(S^k\Omega^k) + \log\det(\Omega^k) + \lambda \left\|\Omega^k\right\|,$$

where $\boldsymbol{U}$ is a membership matrix, and $\{\Theta^l\}$ are irreducible and invertible.

Proposition 1. *The loss function $\mathcal{L}_S$ satisfies the conditions for Theorem 3.1, and thus the clustering accuracy for precision matrix model is guaranteed.*

Proof. First, we introduce some useful notations.

Given the membership $\boldsymbol{U}'$, let $\hat{\Theta}^l(\boldsymbol{U}') = \arg\max_{\boldsymbol{U}',\Theta^l}$. Particularly, for each $l \in [r]$, we have

$$\hat{\Theta}^l(\boldsymbol{U}') = \arg\max_{\Theta} -\sum_{k\in I_l'} \langle S^k, \Theta\rangle + |I_l'|\log\det(\Theta) + \lambda|I_l'|\,\|\Theta\|_1,$$

index l

where $I_l' = \{k : u_{kl}' \neq 0\}$ is the index set for the l-th group based on the membership $\boldsymbol{U}'$. The sample-based loss is defined as

$$F(\boldsymbol{U}') = \mathcal{L}_S(\boldsymbol{U}', \hat{\Theta}^l(\boldsymbol{U}')).$$

Correspondingly, define the population-based loss function as

$$l(\boldsymbol{U}, \Theta^l) = \mathbb{E}_S[\mathcal{L}_S(\boldsymbol{U}, \Theta^l)] = -\sum_{k=1}^{K} \operatorname{tr}(\Sigma^k\Omega^k) + \log\det(\Omega^k) + \lambda\sum_{k=1}^{K}\left\|\Omega^k\right\|_1.$$

~~Given the membership $\boldsymbol{U}'$, let $\tilde{\Theta}^l(\boldsymbol{U}') = \arg\max_{\boldsymbol{U}',\Theta^l}$. Particularly, for each $l \in [r]$, we have~~

$$\tilde{\Theta}^l(\boldsymbol{U}') = \arg\max_{\Theta} -\sum_{k\in I_l'} \langle \Sigma^k, \Theta\rangle + |I_l'|\log\det(\Theta) + \lambda|I_l'|\,\|\Theta\|_1. \tag{1}$$

Then, the population-based loss is defined as

$$G(\boldsymbol{U}') = l(\boldsymbol{U}', \tilde{\Theta}^l(\boldsymbol{U}')).$$

Note that $\hat{\Theta}^l(\boldsymbol{U}')$ and $\tilde{\Theta}^l(\boldsymbol{U}')$ do not have closed forms. But both of them only utilize $|I_l'|$ sample covariance(true covariance) matrices based on the membership.

Next, we verify the functions $F(\cdot)$ and $G(\cdot)$ satisfy the conditions in the Theorem 3.1. Let $\{\boldsymbol{U}, \Theta^l\}$ denote the true membership and precision matrices, and define $\hat{\boldsymbol{U}} = \arg\max_{\boldsymbol{U}} F(\boldsymbol{U})$. We also define the confusion matrix $D = [\![D_{ij}]\!] \in \mathbb{R}^{r\times r}$, where $D_{ij} = \sum_{k=1}^{K} \boldsymbol{I}\{u_{ki} = \hat{u}_{kj} = 1\}$.

1. (Self-consistency) First, we consider the explicit formulas for $G(\hat{\boldsymbol{U}})$ and $G(\boldsymbol{U})$.

$$\begin{aligned} G(\hat{\boldsymbol{U}}) &= l(\hat{\boldsymbol{U}}, \tilde{\Theta}^l(\hat{\boldsymbol{U}})) \\ &= \sum_{l=1}^{r}\left[\sum_{k\in\hat{I}_l} -\langle\Sigma^k, \tilde{\Theta}^l(\hat{\boldsymbol{U}})\rangle + |\hat{I}_l|\log\det(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - \lambda|\hat{I}_l|\left\|\tilde{\Theta}^l(\hat{\boldsymbol{U}})\right\|_1\right] \\ &= \sum_{l=1}^{r}\left[\sum_{a=1}^{r} D_{al}\left(-\langle\Sigma^a, \tilde{\Theta}^l(\hat{\boldsymbol{U}})\rangle + \log\det(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - \lambda\left\|\tilde{\Theta}^l(\hat{\boldsymbol{U}})\right\|_1\right)\right], \end{aligned}$$

and

$$\begin{aligned} G(\boldsymbol{U}) &= l(\boldsymbol{U}, \tilde{\Theta}^l(\boldsymbol{U})) \\ &= \sum_{l=1}^{r}\left[-|I_l|\langle\Sigma^k, \tilde{\Theta}^l(\boldsymbol{U})\rangle + |I_l|\log\det(\tilde{\Theta}^l(\boldsymbol{U})) - \lambda|I_l|\left\|\tilde{\Theta}^l(\boldsymbol{U})\right\|_1\right] \\ &= \sum_{l=1}^{r}\left[\sum_{a=1}^{r} D_{al}\left(-\langle\Sigma^a, \tilde{\Theta}^a(\boldsymbol{U})\rangle + \log\det(\tilde{\Theta}^a(\boldsymbol{U})) - \lambda\left\|\tilde{\Theta}^a(\boldsymbol{U})\right\|\right)\right]. \end{aligned}$$

Define the function

$$h^k(\Theta) = -\langle\Sigma^k, \Theta\rangle + \log\det(\Theta) - \lambda\|\Theta\|_1.$$

By the definition (1), we know that

$$\tilde{\Theta}^k(\boldsymbol{U}) = \arg\max_{\Theta} h^k(\Theta), k = 1, ..., r.$$

Therefore, **we have the self-consistency** of $\boldsymbol{U}$, i.e., $G(\boldsymbol{U}') \le G(\boldsymbol{U})$.

Next, we want to find the function which links the subtraction $G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U})$ with the misclassification rate $MCR(\hat{\boldsymbol{U}}, \boldsymbol{U})$, where $MCR(\hat{\boldsymbol{U}}, \boldsymbol{U}) = \max_{l, a\neq a'\in[r]} \min\{D_{al}, D_{a'l}\}$.

Suppose $MCR(\hat{\boldsymbol{U}}, \boldsymbol{U}) \ge \epsilon$. There exist $l, k \neq k' \in [r]$ such that $\min\{D_{kl}, D_{k'l}\} \ge \epsilon$.

Then, we have

$$\begin{aligned} G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) &\le D_{kl}\left(h^k(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^k(\tilde{\Theta}^k(\boldsymbol{U}))\right) + D_{k'l}\left(h^k(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^k(\tilde{\Theta}^{k'}(\boldsymbol{U}))\right) \\ &\le \epsilon C(\boldsymbol{U}, \Theta^l, \lambda), \end{aligned}$$

where C is a function of the true parameters $\{\boldsymbol{U}, \Theta^l\}$. Need to figure our the explicit form of C in next step.

2. (Bounded difference between sample- and population-based loss) For arbitrary $\boldsymbol{U}$, consider the absolute subtraction

$$\begin{aligned}|F(\boldsymbol{U}) - G(\boldsymbol{U})| &= |\mathcal{L}_S(\boldsymbol{U}, \hat{\Theta}^l(\boldsymbol{U})) - l(\boldsymbol{U}, \tilde{\Theta}^l(\boldsymbol{U}))| \\ &\leq |\mathcal{L}_S(\boldsymbol{U}, \hat{\Theta}^l(\boldsymbol{U})) - l(\boldsymbol{U}, \hat{\Theta}^l(\boldsymbol{U}))| + |l(\boldsymbol{U}, \hat{\Theta}^l(\boldsymbol{U})) - l(\boldsymbol{U}, \tilde{\Theta}^l(\boldsymbol{U}))| \\ &= M_1 + M_2.\end{aligned}$$

Conjecture:

For M_1,

use Cauchy-Schwarz inequality.
multiple ways to bound by, l_1, l_2, l_\infty norm, etc.

$$M_1 = |\sum_{l=1}^{r} \sum_{k \in I_l} \langle (\Sigma^k - S^k), \hat{\Theta}^l(\boldsymbol{U}) \rangle| = \max_{k,(ij)} |\Sigma_{ij}^k - S_{ij}^k| C_1(\boldsymbol{U}, \Theta^l, p),$$

where C_1 is a function of the true parameters $\{\boldsymbol{U}, \Theta^l\}$ and the dimension p.

For M_2, note that $l(\boldsymbol{U}, \Theta)$ is a convex function of Θ and thus l is local Lipschitz. We may have

$$M_2 \leq \max_{l \in [r]} \sup_{\Theta^l} \left| \frac{\partial}{\partial \Theta^l} l(\boldsymbol{U}, \Theta^l) \right| \left\| \hat{\Theta}^l(\boldsymbol{U}) - \tilde{\Theta}^l(\boldsymbol{U}) \right\|_{\max}$$

Also, we can consider $\max_{l \in [r]} \sup_{\Theta^l} \left| \frac{\partial}{\partial \Theta^l} l(\boldsymbol{U}, \Theta^l) \right| = C_2(\boldsymbol{U}, \Theta^l, \lambda)$, where C_2 is the function of the true parameters $\{\boldsymbol{U}, \Theta^l\}$ and tuning parameter λ. Since $\hat{\Theta}^l$ is the sample-based estimation and $\tilde{\Theta}^l$ is the population-based estimation, my conjecture is that $\left\| \hat{\Theta}^l(\boldsymbol{U}) - \tilde{\Theta}^l(\boldsymbol{U}) \right\|_{\max} = C_3(\max_{k,(ij)} |\Sigma_{ij}^k - S_{ij}^k|)$.

Therefore, we bound the difference as

$$|F(\boldsymbol{U}) - G(\boldsymbol{U})| \leq C'(\boldsymbol{U}, \Theta^l, p, \lambda) C''(\max_{k,(ij)} |\Sigma_{ij}^k - S_{ij}^k|),$$

and then we can utilize of residual to find a $p(t) = \mathbb{P}(|F(\boldsymbol{U}) - G(\boldsymbol{U})| \geq t) \to 0$ as $t \to \infty$.

□

2 Misclassification error

We explore the perturbed version of the self-consistency in this section.

Lemma 1 (Self-consistency of $\boldsymbol{U}$)**.** *Suppose $MCR(\hat{\boldsymbol{U}}, \boldsymbol{U}) \geq \epsilon$ and the minimal gap between $\{\Theta^l\}$ denoted δ is positive. For $\lambda \leq C' \left(\frac{\log p}{n}\right)^{1/2}$ with some constant C', we have the perturbation version of the self-consistency.*

how? additional condition on \lambda

$$G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) \leq -\frac{\epsilon}{4\tau^2}\delta^2 + \epsilon\lambda\sqrt{p}C \left(\frac{p \log p}{n}\right)^{1/2} < 0.$$

Proof. Suppose $MCR(\hat{\boldsymbol{U}}, \boldsymbol{U}) \geq \epsilon$. Let $\{\boldsymbol{U}, \Theta^l\}$ denote the true parameters, and $\Theta^l = (\Sigma^l)^{-1}$. Define the function

$$h^k(\Theta) = -\langle \Sigma^k, \Theta \rangle + \log\det(\Theta) - \lambda \|\Theta\|_1 .$$

There exist $l, k \neq k' \in [r]$ such that $\min\{D_{kl}, D_{k'l}\} \geq \epsilon$. Then, we have

$$\begin{aligned} G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) &\leq D_{kl}\left(h^k(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^k(\tilde{\Theta}^k(\boldsymbol{U}))\right) + D_{k'l}\left(h^{k'}(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^{k'}(\tilde{\Theta}^{k'}(\boldsymbol{U}))\right) \qquad (2) \\ &\leq D_{kl}\left(h^k(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^k(\Theta^k)\right) + D_{k'l}\left(h^{k'}(\tilde{\Theta}^l(\hat{\boldsymbol{U}})) - h^{k'}(\Theta^k)\right), \end{aligned}$$

how to control other (k, l, k')

where the second inequality follows the fact that $h^k(\Theta^k) \leq h^k(\tilde{\Theta}^k(\boldsymbol{U}))$ since $h^k(\tilde{\Theta}^k(\boldsymbol{U}))$ is the maximizer of $h^k(\Theta)$ by the definition. For simplicity, let $\hat{\Theta}$ denote $\tilde{\Theta}^l(\hat{\boldsymbol{U}})$. Define $\Delta^k = \hat{\Theta} - \Theta^k$. Consider the function

$$f^k(t) = \log\det(\Theta^k + t\Delta),$$

and by Taylor expansion we have

$$f^k(1) - f^{k'}(0) = \langle \Sigma^k, \Delta^k \rangle - \mathrm{vec}(\Delta^k)^T \int_0^1 (1-v)(\Theta^k + v\Delta^k)^{-1} \otimes (\Theta^k + v\Delta^k)^{-1} dv \mathrm{vec}(\Delta^k).$$

Then, we have

$$\begin{aligned} h^k(\tilde{\Theta}^k) - h^k(\hat{\Theta}^k) &= \langle \Sigma^k, \Delta^k \rangle - f^k(1) + f^k(0) - \lambda\left(\left\|\Theta^k\right\|_1 - \left\|\hat{\Theta}\right\|_1\right) \\ &\geq A_1 - |A_2|, \end{aligned}$$

where

$$\begin{aligned} A_1 &= \mathrm{vec}(\Delta^k)^T \int_0^1 (1-v)(\Theta^k + v\Delta^k)^{-1} \otimes (\Theta^k + v\Delta^k)^{-1} dv \mathrm{vec}(\Delta^k) \\ A_2 &= \lambda\left(\left\|\Theta^k\right\|_1 - \left\|\hat{\Theta}\right\|_1\right). \end{aligned}$$

By Guo's paper, we know that

$$A_1 \geq \frac{1}{4\tau^2}\left\|\Delta^k\right\|_F^2, \qquad (3)$$

where $\max_{k\in[r]} \varphi_{\max}(\Theta^k) \leq \tau < \infty$. Also note that

$$|A_2| \leq \lambda\left\|\Theta^k - \hat{\Theta}\right\|_1 \leq \lambda\sqrt{p}\left\|\Delta^k\right\|_F. \qquad (4)$$

Plug the inequalities (3) and (4) in to the inequality (2), we obtain that

$$G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) \leq D_{kl}\left(-\frac{1}{4\tau^2}\left\|\Delta^k\right\|_F^2 + \lambda\sqrt{p}\left\|\Delta^k\right\|_F\right) + D_{k'l}\left(-\frac{1}{4\tau^2}\left\|\Delta^{k'}\right\|_F^2 + \lambda\sqrt{p}\left\|\Delta^{k'}\right\|_F\right).$$

Intuitively, if we have λ very small, then we obtain the perturbation version of self-consistency. By a straightforward calculation, if we have

$$\lambda \leq \frac{1}{4\tau^2\sqrt{p}} \min_{k\in[r]} \left\|\Delta^k\right\|_F, \qquad (5)$$

then the perturbation version of self-consistency holds. Recall our previous conclusion for the Ω estimation. If $\lambda = \mathcal{O}\left(\left(\frac{\log p}{n}\right)^{1/2}\right)$, we have

not directly applicable.
Your earlier result (0115.pdf) for Omega is unc
constrained optimization, but not under
penalized optimization. (?)

$$\min_{k\in[r]} \left\|\Delta^k\right\|_F \leq C\left(\frac{p\log p}{n}\right)^{1/2}.$$

Extension should be easy though.

with high probability. This implies that when $\lambda \leq C'\left(\frac{\log p}{n}\right)^{1/2}$, the λ satisfies the condition (5) with high probability. Finally, we obtain the perturbation version of self-consistency,

$$\begin{aligned} G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) &\leq -\frac{\epsilon}{4\tau^2}\left\|\Theta^k - \Theta^{k'}\right\|_F^2 + \epsilon\lambda\sqrt{p}C\left(\frac{p\log p}{n}\right)^{1/2} \qquad \text{how?} \\ &\leq -\frac{\epsilon}{4\tau^2}\delta^2 + \epsilon\lambda\sqrt{p}C\left(\frac{p\log p}{n}\right)^{1/2}, \end{aligned}$$

where δ is the minimal gap between Θ^l. □

Remark 1. When $\lambda = 0$, the subtraction $G(\hat{\boldsymbol{U}}) - G(\boldsymbol{U}) \leq -\frac{1}{4\tau^2}\delta^2$ agrees with the result under the case without penalty.

Remark 2. The difficulty of the proof comes from that $\tilde{\Theta}^l(\boldsymbol{U})$ does not have a closed form. In other literatures, they usually consider the true Θ^l rather than $\tilde{\Theta}^l(\boldsymbol{U})$ under the true membership. The possible reason is that the properties (such as singular value, minimal gap) of Θ^l are easy to describe while it is hard to tell the properties of $\tilde{\Theta}^l(\boldsymbol{U})$ (except it is an optimizer). Therefore, I introduce the true precision matrices in the proof in step (2). As a result, the upper bound becomes related with the precision matrices estimation $\|\Delta\|_F = \left\|\hat{\Theta} - \Theta^k\right\|_F$, and thus the control for λ is required.

3 Others

Theorem 3.1 (General property for loss function to guarantee the clustering accuracy)**.** *Let $\{\mathcal{C}, \boldsymbol{M}_k\}$ denote the true parameters, and $\mathcal{L}_{\mathcal{Y}}(\mathcal{C}', \boldsymbol{M}_k')$ denote the sample-based loss function. Define the sample-based loss function with respect to $\boldsymbol{M}_k'$ as*

$$F(\boldsymbol{M}_k') = \mathcal{L}_{\mathcal{Y}}(\hat{\mathcal{C}}(\boldsymbol{M}_k'), \boldsymbol{M}_k'),$$

where

$$\hat{\mathcal{C}}(\boldsymbol{M}_k') = \arg\max_{\mathcal{C}} \mathcal{L}_{\mathcal{Y}}(\mathcal{C}, \boldsymbol{M}_k').$$

Correspondingly, define the population-based loss function with respect to $\boldsymbol{M}_k'$ as

$$G(\boldsymbol{M}_k') = l(\tilde{\mathcal{C}}(\boldsymbol{M}_k'), \boldsymbol{M}_k'),$$

where

$$l(\mathcal{C}, \boldsymbol{M}_k) = \mathbb{E}_{\mathcal{Y}}[\mathcal{L}_{\mathcal{Y}}(\mathcal{C}, \boldsymbol{M}_k)], \quad \textit{and} \quad \tilde{\mathcal{C}}(\boldsymbol{M}_k') = \arg\max_{\mathcal{C}} l(\mathcal{C}, \boldsymbol{M}_k').$$

Suppose the loss function satisfies the following properties

1. *(Self-consistency to $\boldsymbol{M}_k$) Suppose $MCR(\boldsymbol{M}_k', \boldsymbol{M}_k) \geq \epsilon$ for $\epsilon > 0$. We have*

   $$G(\boldsymbol{M}_k') - G(\boldsymbol{M}_k) \leq -C(\epsilon),$$

   where $C(\cdot)$ takes positive values.

2. *(Bounded difference between sample- and population-based loss) The difference between sample-based and population-based loss function is bounded in probability, i.e.,*

   $$p(t) = \mathbb{P}(|F(\boldsymbol{M}_k') - G(\boldsymbol{M}_k')| \geq t) \to 0, \quad \textit{as} \quad t \to \infty.$$

Let $\{\hat{\boldsymbol{M}}_k\}$ *be the maximizer of* $F(\boldsymbol{M}_k)$*. Then, we have the following clustering accuracy, for any* $\epsilon > 0$*,*

$$\mathbb{P}(MCR(\hat{\boldsymbol{M}}_k, \boldsymbol{M}_k) \geq \epsilon) \leq p\left(\frac{C(\epsilon)}{2}\right).$$