

Jun 15 - Jun 22 Seminar Review

Jiaxin Hu

June 16, 2020

1 MDFS 6.16

Title: Representation, Modeling, and Gradient Based Optimization in Reinforcement Learning

Representer: Sham Kakade

How can we address the large/infinite table in Reinforcement Learning? This seminar discusses this problem from three aspects: representation, modeling, and optimization. In representation, the log-linear policy method with 100 parameters is shown to perform no worse than a full neural network with 1000 parameters. That implies that fewer features also support good learning. A hard threshold to the goodness of the features is also provided. In modeling, Sham discusses the sampling efficiency and provides a minimax optimal sampling criterion to achieve a good performance. In optimization, training from different configurations is suggested to strengthen the robustness of the model. The Natural Policy Gradient (NPG) is proposed to solve the non-convex optimization. The global convergence of NPG is independent of dimension. The theoretical performance of NPG for the linear policy method is also guaranteed.